

File System Naming and Reliability

Naming

- File needs to have a handle that allows for references
- OS uses simple numbers as names, which aren't usable by people
- Need a better way to name files which is user friendly and allows for easy organization

Hierarchical Name Spaces

- Graphical organization
- Typically organized using directories
 - A file containing references to other files
- Nested directories can form a tree
 - A file name is a path through the tree
 - Can form a directed graph if files are allowed to have multiple names

Directories are Files

- Directories are a special type of file that maps file names to associated files
- Directory contains multiple directory entries
 - Each entry describes open file and its name
- User apps are allowed to read directories
 - Can get info about files and allowed to see which file exist
 - Usually write permissions are granted by the OS

Traversing the Tree

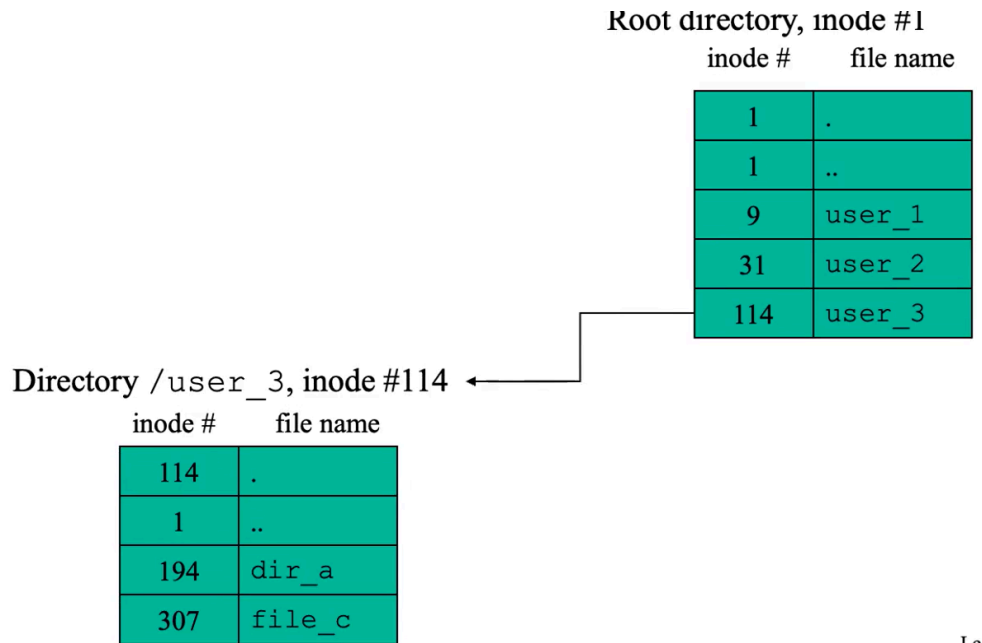
- Entries in directories point to child directories
 - Describes a lower level in the hierarchy
- To name a file at that level, name the parent directory and the child directory
- Moving up the hierarchy is often useful
 - Directories have special entries to point to the parent “..”

File Names vs Path Names

- True Names
 - One possible name for the file
 - In DOS, file is described by a directory entry

Hard Links

- File names separated by slashes
 - /user/fdsf
- Directory entries only point to inodes and are a mapping between file names and inodes
- Hard Links / Directory entry
 - File name and inode number, known as “hardlink”
 - Multiple directory entries point to the same inode



Let

Multiple File Names in Unix

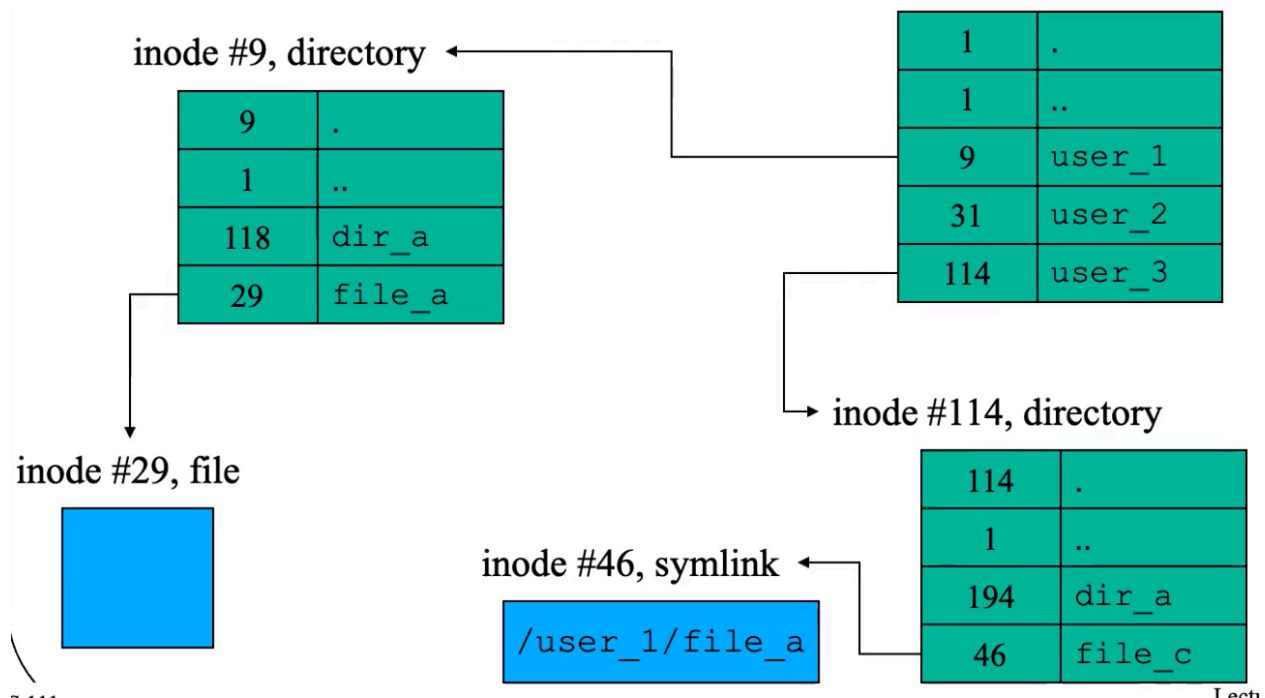
- How links relate to files
 - They're only names
- All other metadata is stored in the file inode
- All links provide the same access to the file
 - Anyone with read access can create a new link
- All links are equal
 - Nothing special about the first link

Links and Deallocation

- If files exist under many names, can't remove the data until all names have been removed
- Hard link count is stored in the inode

Symbolic Links

- Different way of giving files multiple names
- Symbolic links implemented as a special type of file
 - Indirect reference to some other file
 - Contents is a path to another file
- File system recognizes symbolic links
- Symbolic link isn't a reference to the inode
 - Symbolic links don't prevent deletion
 - Deleting the symbolic name doesn't change the link count



File Systems Reliability

- Multiple operations
 - Writing one or more metadata blocks
 - Inode, free list, and directory blocks
- All must be committed to disk for the write to succeed
- Each block write is a separate hardware operation

Deferred Writes

- Process allocates a new block to file A
 - Get a new block from the free list
 - Write out the updated inode for file A
 - Defer free list write back
- System crashes and reboot
 - New process wants a new block for file B
 - But file A is already using the block to store data

Core Reliability Problem

- File system writes typically involve multiple operations
 - Not just writing a data block to disk/flash
 - Writing one or more metadata blocks
- All must be committed to disk for the write to succeed
 - Failure on one operation could be bad
- Each block write is a separate hardware operation

Deferred Writes

- Process allocates a new block to file A
 - We get a new block from the free list
 - We write out the updated inode for file A
 - Defer free list write back (delay the write operation to batch it with more data)
- System crashes and after it reboots
 - New process wants a new block for file B
 - We get a block from the stale free list
- Two different files now contain the same block
 - File A is written, causing B to get corrupted and vice versa

App expectations When Writing

- App makes sys call to perform writes
- Upon syscall return, we expect the write to be safe
- We can block the writing app until its safe

Buffered Writes

- Don't wait for the write to actually be persisted
- Keep track of it in RAM
- Tell the app its OK
- At some later point actually write to persistent memory
- Advantages
 - Less app blocking
 - Deeper and optimizable write queries
- Disadvantages
 - Crashes could cause problems with memory allocation

Ordered Writes

- Carefully ordered writes can reduce potential damage
- Write out data before writing pointers to it
 - We want to make sure the data exists before creating the pointer to it
 - Unreferenced objects can be garbage collected
 - Pointers to incorrect info can be serious
- Write out deallocations before allocations
 - Disassociate resources from old file immediately
 - Free list can be corrected by garbage collection
 - Improperly shared data is more serious than missing data (better to lose data than to corrupt the file system)

Practicality of Ordered Writes

- Greatly reduced I/O performance
 - Eliminates accumulation of near by operations
 - Eliminated consolidation of updates to same block
- May not be possible due to device constraints
- Doesn't solve the problem
 - Ordered writes don't eliminate incomplete writes

Audit and Repair


- Design file system structures for audit and repair
- Audit file system for correctness and use redundant info to enable automatic repair
- Used to be standard practice, but isn't practical since its so slow

Journaling

- Part of our storage is devoted to a journal, which keeps track of everything that I intend to write
- Circular buffer journaling device
 - Always sequential
 - Writes can be batched
 - Journal is relatively small
- Efficiently schedule actual file system updates
 - Journal keeps track of the ordering that updates need to be made in
- Journal completions when real writes happen
 - Once journal updates are made, the contents that have been written don't need to be stored anymore
- Upon system crash, go to the journal to see what needs to be done
 - Even if the journal wasn't cleared during the crash or there was an intermediate step crash, it will still work

Journal example

- Write 2 pages
 1. Put start record into the journal
 2. Put the 2 pages into the journal
 3. Put the end record into the journal
 4. Overwrite the data, then delete entries from the journal



Batching Journal Entries

- Operation is safe after journal entry persisted
- Small writes are still inefficient
- Accumulate batch until full or max wait time

Journal Recovery


- It's a small circular buffer that can be recycled after old ops have completed
 - Time stamps distinguish new entries from old
- After the system restarts
 - Review entire journal
 - Note which ops were known to have completed
 - Perform all writes not known to have completed
- Truncate journal and resume normal operation

Journalling Functionality

- Journal writes are much faster than data writes
- In normal operation, journal is write only
 - File system never reads/processes the journal
- Scanning the journal on restart is fast
 - Very small
 - Can read sequentially with large and efficient reads
 - All recovery processing is done in memory
- Journal pages may contain information for multiple files
 - Performed by different processes and users

Metadata Only Journaling

- Journal meta-data
 - Only contains the metadata
 - Small and random
 - Integrity critical
- Journal data
 - Large and sequential
 - Less order sensitive
- Safe metadata journaling
 1. Allocate new space for the data and write it there
 - a. Data could take a lot of space, so we should write it to disk instead of storing in journal
 2. Journal the metadata updates
 - a. Consistency in the metadata is extremely important
 - b. Metadata points to the data that we stored on disk



Log Structured File Systems

- Journal is the file system (called a log)
 - All inodes and data updates written to the log
 - All updates in a log are updates using **redirect on write** (don't overwrite the old data, write it elsewhere then change the metadata pointer to it)
 - This means that old versions remain, but the inode pointer points to new data
 - In memory index caches inode locations

- Dominant architecture
 - Flash file systems
 - Key/value store
- Issues
 - Recovery time (to reconstruct index/cache)
 - Log defragmentation and garbage collection

Navigating a Logging File System

- Inodes point at data segments (pages) in the log
 - Sequential writes may be contiguous in log
 - Random updates can be spread across the log
- Updated inodes are added to the end of the log
- Index points to latest version of each inode
 - Index is periodically appended to the log
- Recovery
 - Find and recover the latest index
 - Replay all log updates since then

Redirect on Write

- Once written, blocks and inode are immutable
 - Add new info to the log and update the index
- Old inodes and data remain in the log
 - If we have an old index, we can access them
- Garbage collections is used to recycle old log entries