

## OS Interfaces and Abstractions

### API

- Source level interface
  - Includes files, data types, constants, macros, routines, and parameters
- Basis of software portability
  - Recompile program for the desired architecture
  - Resulting binary runs on that architecture and OS
- An API compliant program will compile and run on any compliant system
  - APIs are primarily for programmers
  - Should work on all types of machines

### ABI

- Binding interface for specifying:
  - Dynamic loadable libraries (DLLs)
  - Data formats, calling sequences, linkage conventions
- Binding interface of an API to a hardware
- Basis for binary compatibility
  - Usually one ABI for an OS
  - One binary servers all customers for that hardware (all x86, linux, MacOS)
- Installing a new version of the OS should allow your current apps to still work
- ABI compliant program will run on any compliant system

### Libraries and Interfaces

- Normal libraries are accessed through an API
  - Source-level definitions of how to access the library
  - Readily portable between different machines
- Dynamically loadable libraries also called through an API
  - But the dynamic loading mechanism is ABI specific
  - Issues of word length, stack format, linkages

### Interfaces and Interoperability

- Strong, stable interfaces are key to allowing programs to operate together
- Also important to OS evolution, since we don't want to break existing programs
- Interface (API and ABI) between the OS and apps shouldn't change

## Interface Interoperability

- Stability
  - Programs use sys calls, libraries, and external files
  - API requirements are frozen at compile time
  - Execution platform must support those interfaces
  - Future upgrades must support old interfaces
- Compliance
  - Comply with the interface, don't try to find loopholes to interact with hardware
  - Complete interoperability testing is impossible (cannot test all apps on all platforms)
  - New apps are continuously added, which will comply with the interface

### Side Effects

- Side effect: occurs when an action has an unexpected outcome
- Effects not specified by interfaces and aren't intended
- Focus on the interface, don't try to exploit side effects

### Abstractions

- Simplifies the complex implementations for interacting with hardware
  - Hides error handling, performance optimization
  - Eliminates behaviors not important to the programmer
1. Memory abstractions
  2. Processor abstractions
  3. Communications abstractions

## Memory Abstractions

- Variables, Files, Database records, messages

### Complications

- Persistent memory: memory which persists after a computer is shutdown
- Transient Memory:
- Coherence and atomicity
- Latency (reading or writing to flash drive is very slow, writes takes different amounts of time)
- Same abstractions can be implemented in differently in different devices (read only, solid state, etc)

### Source of complications

- OS doesn't have abstract devices
- Physical devices specific to that machine
- We must abstract the hardware regardless of what the implementation looks like

## Memory Example

- File, can read or write to the file
- Coherence: If we write the file, we expect our next read to reflect the results of the write
- Atomicity: We expect the entire read/write to occur, entire chunk of 500 bytes should be written instead of just a portion
- If there are several read/writes to the files, we expect them to occur in some order

## Implementation flash drives

- Flash drive
- Write once semantics:
  - Rewriting requires an erase cycle
  - Erases a whole block
  - Erases are slow
- Atomicity occurs at the work level
- Blocks can only be erased so many times
- OS needs to abstract the above behaviors

## Abstraction considerations

- Flash drive storage is split into blocks
- Different structures for the file system since it's difficult to overwrite data words in place
- Garbage collection to deal with blocks filled with inactive data
- Maintaining a pool of empty blocks (we can quickly write to these blocks)
- Wear leveling in use of blocks (make sure blocks are evenly spaced)
- Something that provides atomicity to multiple writes

## Interpreter Abstractions

- An interpreter is code that performs commands
- **Repertoire**: a set of things the interpreter can do (can the CPU do floating point operations?)
- **Environment reference**: Describes the current state on which the next instruction should be performed
- **Interrupts**: Out of order instruction execution or canceling execution

## Process Abstraction

- Every program is abstracted to a process
- Repertoire : Source code defines what the code can do
- Environment: Stack, heap, register contents
- Interrupts: Ctrl - C

## Process Abstractions

- Simple with one process, but difficult with multiple
- OS has limited physical memory to hold the environment information
- There is usually one set of registers
- No other interpreters should interfere with a process's resources

### Abstraction Implementation

- Schedulers to share the CPU among various processes
- Memory management hardware and software
  - To multiplex memory use among the processes
  - Gives the impression of full use of memory (even though memory is actually shared by different processes)
- Access control mechanisms for other memory abstractions
  - So other processes can't read or alter my files
  - OS needs to isolate processes

## Communications Abstractions

- Communication link allows one interpreter to talk to another on the same or different machines
- Memory and cables or networks and interprocess communication mechanisms

### Communication properties

- Highly variable performance (network latency can be slow or packets lost)
- Often async, resulting in issues with syncing both parties
- Receiver may only perform the operation since the send occurred
- Additional complications when working with a remote machine

### Same Machine Process Communication

- Easy if both processes are on the same machine
  - Done by reading and writing to RAM
  - Or transfer control of memory containing the memory from the sender to receiver

### Different Machine Communication

- Need to optimize costs of copying, resulting in difficult memory management
- Inclusion of complex network protocols in the OS itself
- Difficulties with message loss and retransmission
- Security concerns

## Generalizing Abstractions

- Different things are made to appear the same
- Applications can all deal with a single class, or a common unifying model

- Portable Document Format (PDF) for printed output
- SCSI/SATA for disks

### Federation Frameworks

- A generalized abstraction
- A structure that allows many similar but different things to be treated uniformly
- By creating one interface that works with all implementations (at a high level, users can use flash drives and hard drives in the same way)
- Least common denominator which works for multiple models:
  - Common model has optional features
  - Ex. printers: some printers can print double sided, others can't

### Abstractions and Layering

- Common to create increasingly complex services by layering abstractions
  - Generic file system layers on a particular file system
- Layering allows good modularity
  - Easy to build multiple services on a lower layer
  - Multiple file systems on one disk

#### Downside of Layering

- More abstractions results in performance penalties
- Often expensive to move from one layer to another
  - Since frequently requires changing data structures or representations
  - Abstractions require additional instructions
- Another downside is that lower layer may limit what the upper layer can do
  - Any abstract network link may hide causes of packet losses

#### Other OS Abstractions

- Different abstractions provide different types of abstractions
- OS must do work to provide an abstraction (higher level = more work)
- Programmers and users have to choose the right abstraction to work with

### Summary

- Stable interfaces needed for proper performance of an OS
  - APIs for program development and ABIs for user experience
- Abstractions make OS systems easier for both programmers and users
- The most important OS abstractions involve memory, interpreters, and communication