

# Scheduling

## OS Scheduling choices

- What job to run next
- How long to let the process run
- What order to handle a set of block requests
- If multiple messages to be sent over the network, what order should they be sent

## Scheduling Goals

- Design a scheduling algorithm to achieve goals
- Different algorithms try to optimize different metrics
- Changing algorithm changes system behaviors
- Policy decision

## Process Queue

- OS keeps a queue of processes that are ready to run
  - Ordered by the task to run next
- Processes that aren't ready to run:
  - Aren't in the queue
  - At the end of the queue
  - Ignored by the scheduler

## Potential Goals

- Maximize throughput
  - Maximize user work done
  - Running OS code doesn't count
- Minimize the average waiting time
  - Avoid delaying too many tasks
- Ensure fairness
  - Minimize worst case waiting time
  - Attempt to distribute the work being done among all tasks
- Meet priority goals
  - Certain tasks have a higher priority than others and should be done first
- Real time scheduling
  - Scheduled items tagged with a deadline that must be met

## Different Kinds of Systems

- Time Sharing
  - Fair response time to interactive programs

- Each user gets an equal share of the CPU
- Ex. Many users running programs on a server
- Batch
  - Maximize total system throughput
  - Delays of individual processes is unimportant
  - Ex. Supercomputers, we want to maximize the amount of work we finish efficiently
- Real Time
  - Critical operations must happen on time
  - Ex. playing music, streaming
- Service Level Agreement
  - To share resources between multiple customers
  - Agreement which promises a certain amount of resources to the customer
  - Cloud providers must try to rent out all of their services while maintaining the promises, efficient usage is done by scheduling

#### Scheduling: Policy and Mechanism

- **Dispatching: Process in which the scheduler moves jobs onto and off the processor core**
- Dispatching is irrelevant of the scheduling algorithm
- Dispatching shouldn't depend on the policy to decide who to dispatch
- Separate policy and the dispatching mechanism

#### Preemptive v.s Non preemptive Scheduling

- When we schedule a piece of work:
  - We could let it use the resource until it finishes
  - Or interrupt it to work on another process
- **Non preemptive:** scheduled work always runs to completion
- **Preemptive:** Scheduler halts a task to run another task

#### Non Preemptive Scheduling

- Advantage
  - Low scheduling overhead since we always run the process until it's finished
  - **Produces high throughput**
  - Conceptually simple and easy to implement
- Disadvantage
  - Poor response time

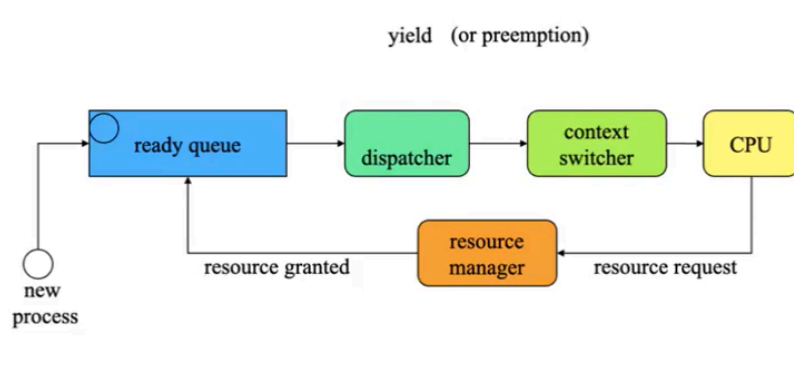
- Bugs can freeze the machine or cause infinite loop, since OS doesn't interrupt the process
- Poor fairness, since a long process can keep running
- Make realtime and priority scheduling difficult

## Preemptive Scheduling

- Advantages
  - Good response time
  - Very fair usage
  - Good for real time and priority scheduling
- Disadvantages
  - More complex
  - Requires ability to halt process and save state
  - Throughput suffers
  - Higher overhead

## Scheduling

- Yield: When a process gives up its spot and allows another process to run on the CPU




## Quantifying metrics

- Throughput
  - (finished process / unit of time)
  - Different processes need different run times
  - Completion time isn't controlled by scheduler
- Delay
  - Not clear hat delays measure (time to finish job, time to get response)


### No Ideal Throughput

- Scheduling is costly
  - It takes time to dispatch a process
  - More dispatches means more overhead
  - Less time available to run processes
- Minimizing the performance gap
  - Reduce the overhead per dispatch
  - Minimize the number of dispatches (per second)



### No Ideal Response Time

- Real systems have finite limits
  - Such as queue size
- When limits exceed, request are typically dropped
  - Results in infinite response time
  - May be automatic retries, which could also get dropped
- Too much load
  - Lots of dropped items
  - During heavy load, overheads will increase drastically



## Graceful Degradation

- System overloaded
  - When it can no longer meet service goals
  - Continued service results in degraded performance
  - Maintains performance by rejecting work
  - Resume normal service when load returns to normal
- Things NOT to do when overloaded
  - Allow throughput to drop to zero
  - Allow response time to grow without limit

## Non Preemptive Scheduling

- Scheduled process continues running until it yields CPU
- Works well for simple systems
  - Small number of processes
  - With natural producer consumer relationships
- Good for maximizing throughput
- Spends on each process to voluntarily yield
  - A buggy process which never yields will freeze the whole system

## First Come First Serve

- Run first process in the ready queue until it completes/yields
- Continue the next process in the queue
- Results high variable delays and poor response times, since a long process will run to completion

### Use Cases

- Response time not important
- Minimizing overhead is more important than any single job's completion time
- Ex. Embedded systems when computations are brief

## Real Time Schedulers

- Some things must happen in particular time
- In a video, one frame must come after another
- Systems must schedule based on real time deadlines

## Hard Real Time Schedulers

- **Hard deadlines**: system is configured to absolutely meet the deadline (controlling nuclear power plant)
- Schedule is worked out ahead of time

### Ensuring hard deadlines

- Deep understanding of the code and how long it takes to run
- Avoid non deterministic timings (every action is predictable)

- Turn off interrupts (which is non deterministic)
- Scheduler is non preemptive (everything is planned beforehand)
- Runs on a predefined schedule (no real time decisions are made)

### Soft Real Time Schedulers

- Goal of scheduler is to avoid missing deadlines (although some can be missed)
- Not as vital to run tasks to completion
- Does not need as much analysis
- Missing a deadline could cause system to fall behind or drop future job

#### Algorithms

- Earliest Deadline First
- Job has a deadline and queue is sorted based on earliest deadline
- Always pick the job with the earliest deadline
- Goal: minimize total lateness

### Preemptive Scheduling

- A running process can continue running, yields, or OS can interpret it
- An interrupt allows another process to run and the interpreted process can restart later
- A process can be forced to yield
  - If there is a more important process (such as result of an I/O completion)
  - If the running processes importance decreased, it can be replaced by a higher priority one
- Interrupted process might not be in a clean state
  - Complicated saving and restoring its state
- Fluid context switches needed to avoid resource sharing problems
  - More context switches = higher overhead

### Implementing Preemption

1. Need OS to regain control from the process
  - Syscall or clock interrupt
2. Consult scheduler
  - Has any process increased its priority
  - Has any process been awakened
  - Has current process had its priority lowered
3. Scheduler finds highest priority ready process
  - If current process is highest, then continue running
  - If not, replace it with a higher priority process

### Clock interrupts

- Can generate an interrupt at a fixed time interval to temporarily halt a running process
- Prevents runaway infinite process so it doesn't keep control forever
- Important for preemptive

### Round Robin Algorithm

- Fair share scheduling
  - All processes given equal amounts of time to run
  - All processes have the same delays in the queue
- All processes re assigned a normal time slice
  - Usually the same sized slice for all
  - Each process is scheduled in these time slices, then placed back into the queue

### Advantages

- All processes get relatively quick chance to do some computation
  - Good for interactive processes (allows for scheduling tasks between waiting for I/O)
  - More responsive: All processes have a chance to execute relatively quickly
- Runaway processes do relatively little harm since they are stopped by interrupts

### Disadvantages

- No process is quickly finished
- Far more context switching, extra overhead

### Choosing Time Slice

- Performance of a preemptive scheduler depends heavily on how long the time slices
- Long time slices avoid context switches
  - Better throughput
- Short time slices and more frequent switches
  - Better response time to process

### Costs of Context Switch

- Entering the OS requires saving registers and calling scheduler
- Cycles to choose who to run
- Moving OS context to the new process (switching stack and process descriptor)
- Switching process address spaces
- Losing instructions and data caches results in slower next instructions
  - Losing process caches reduces performance

### MLFQ: Priority Scheduling

- **Starvation**: low priority algorithms might not be able to run
  - Solution adjust priority over time
    - Processes that have run for a long time have priority temporarily lowered
    - Processes that haven't run in a long time have increased priority
1. All processes start in the highest priority queue
  2. Move it to a lower priority when it uses its allocated resources or time
  3. Periodically move all processes to a higher priority queue

#### Advantages

- Acceptable response time for interactive jobs
  - Jobs don't have to wait for long before being scheduled
- Efficient but fair CPU use for non interactive job
  - Only run for an allocated time slice
- Automatic adjustment of schedules is done based on the behavior of jobs