


Deadlock

Deadlock

- Each resource holds its lock and does not release it since it needs a lock that someone else holds it



- Problem was carefully designed to cause deadlocks

Deadlock Issues

- Major problem for parallel interpreters
 - Relatively common in complex apps and results in large failures
- Difficult to find through debugging
 - Happen intermittently and are much easier to prevent at design time
- Process resource needs continue to change
 - Depends on what data they are operating on
 - Depends on where in computation they are at and what errors have occurred
- Modern software has many services
 - Many services are not aware of one another
 - Each service has complexity that we aren't aware of (resource usage or how they are serialized)

Resource Types

- Commodity Resources
 - Clients need an amount of a certain resource (memory)
 - Deadlocks result from over commitment
 - Avoidance done in resource manager
- General Resource
 - Clients need a very specific instance of something
 - Ex. Particular file

Basic Deadlock Conditions

- For deadlock, one of the 4 can occur:

Mutual Exclusion

- Resource can only be used by one entity at a time
- If multiple entities can use a resource, then it can be given to all

Incremental Allocation

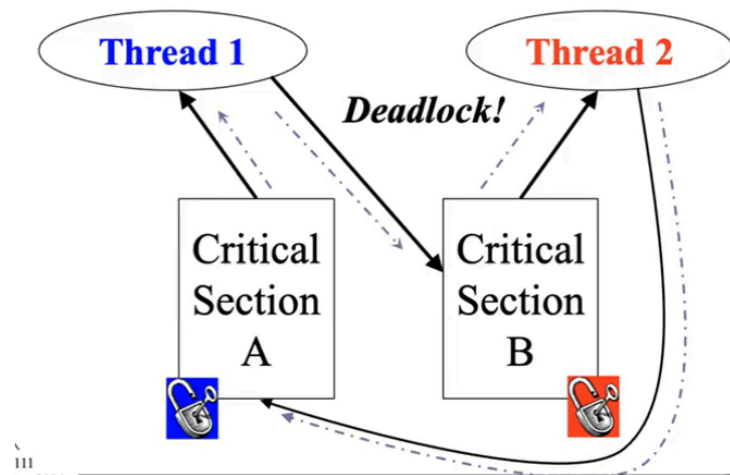
- Processes are allowed to ask for resources whenever they want
- However, if all locks were asked for at once, then no deadlock occurs

No Pre-emption

- When an entity has reserved a resource, can't take it away
- Deadlocks can be resolved by taking away resources

Circular Waiting

- Each process waits on a resource that someone else holds
- Wait for graph:



- 2 Threads, 2 critical sections
- When thread A gets a lock for section A, draw arrow to thread 1 from section A
- When thread B gets a lock for section B, draw an arrow to thread 2
- If thread 1 tries to get the lock for B, then draw arrow from thread 1 to section B (indicates waiting for graph B)
- If thread 2 tries to get lock for A, then draw arrow from 2 to A
- This forms a cycle

Deadlock Avoidance

- Use methods that ensure that no deadlock can occur
- Use advance reservations where no object can over reserve resources

Using reservations

- Advance reservations for commodity resources
 - Overbooking is ok for commodity resources
 - Only grant reservations if resources available
- Over subscriptions are detected early before processes ever receive resources

Overbooking vs Under Utilization

- Processes cannot perfect predict resource needs ahead of time
- Ensure they have enough (tend to ask for more than needed)
- Either the OS
 - Grants requests until everything is reserved
 - Continue granting requests beyond available amount (since not all resources reserved are used)

Handling Reservation Failures

- Reservation eliminates deadlock
- Apps must handle reservation failures
 - App design should handle failures gracefully
 - App must have a way of reporting failure to requester
 - App must continue running (all critical resources were reserved at start up time)

Rejecting Requests

- Better to inform the process they can't have a resource than to promise it when it's been reserved for something else
- Given advance notice, app may be able to adjust service

Deadlock Prevention

- Commodity resource locking doesn't work on general resource locking where there is a general lock on some resource
- **Deadlock avoidance tries to ensure no lock can cause a deadlock**
- **Deadlock prevention tries to ensure that a particular lock doesn't cause a deadlock**
- Targets the four necessary conditions for deadlock, and if any one of these conditions doesn't occur then deadlock can't occur

Mutual Exclusion

- Deadlock requires mutual exclusion, where P1 has the resource before P2 gets it
- Can't deadlock on a shareable resource
 - Use atomic instructions

- Raider writer locking can help (readers can share, writers handled in other ways)
- Can't deadlock process's private resource (giving each process private resource not always practical)

Incremental Allocation

- Deadlock requires you to block holding resources while you continue to ask for more locks
1. Allocate all resources one operation
 - a. Must get all locks at once, otherwise failure
 - b. All or nothing
 2. Non Blocking requests
 - a. A request that can't satisfied immediately will fail
 3. Disallow blocking while holding resources
 - a. Must release all held locks prior to blocking (can't hold locks while spinning)
 - b. Reacquire them after you return to running the process
 - c. Doesn't solve the problem when you return, the locks you need might not be free

No Pre-emption

- Deadlock can be broken by resource confiscation (steal locks)
 - Resource leases with time outs and lock breaking
 - **Lease**: limited time access to a resource (results in issue of partially completed tasks)
 - **Lock breaking**: Resource can be seized and reallocated
- Revocation must be enforced
 - Invalidate previous owners resource handle
 - Reset the partially completed process after its lock is taken
 - If revocation not possible, kill previous owner
- Some resources may be damaged by lock breaking
 - Previous owner was in the middle of critical section
 - May need mechanisms to repair resource
- Resources are designed with revocation
 - If a process crashes, one option is to kill the rest

OS seizing resources

- Process has to use a system service to access the resource, then the service can stop requests
- OS can't revoke a process access if the process has direct access
 - Object is part of address space
 - If the resource isn't controlled by the OS, OS can't handle it

Circular Dependencies

- Total resource ordering
 - All requesters allocate resources in same order
 - Allocate R1 and then R2 after
 - Someone else may have R2 but doesn't need R1
- Assumes we know the exact order of resources
 - Order by resource type (groups before members)
 - Order by relationship (parents before children)
- Lock dance: similar to round robin
 - Release R2, allocate R1, require R2

Lock Dance:

- Must acquire lock on list head
- Individual buffers also have locks
- In order to get the lock for a buffer, the process needs the lock on the list head

Approaches to solving deadlock

- No one universal solution
- Solve each individual problem in any way
 - Making resources shareable when possible
 - Use reservations for commodity resources
 - Ordered locking or no hold and block where possible
 - Last resort: leasing and lock breaking

Ignoring Deadlock

- If it doesn't occur much, we don't implement any anti-deadlock measure

Implement Deadlock Detection

- Identify all resources that can be locked
 - Not always clear in an OS
 - Hard to determine for app level locks
- Deadlocks outside OS difficult to detect
- Other synchronization problems such as bugs in app code can cause issues

Health Monitoring

- Deadlock detection rarely makes sense since its hard to implement and difficult to fix even when detected
- Service/app level health monitoring is better
 - Monitor app progress
 - If response takes too long, service is marked as “hung”
- Health mentoring easy to implement
- Can detect a variety of problems
 - Deadlocks, live locks, etc

Live-lock

- Process is running but won't free R1 until it gets message
- Process will send the message is blocked for R1
- Process blocked and waits for completion that never happens

Priority Inversion

- Higher level process holding lock can't run due to long running lower priority process

Monitoring Process Health

- Look for process exists or core dumps
- Passive observation
 - Process consuming CPU time or is blocked
 - PProcess doing network or disk
- External monitoring
 - Test for responses from a system (if a response isn;t received, problem occurred)
 - Pings, null requests, standard test requests
- Internal instrumentation
 - White box audits and monitoring

Handling Unhealthy Processes

- Kill and restart “all of the affected software”
- Kill as few processes as possible
- Apps designed to react to kills and restarts
- Highly available systems define restart groups
 - Group of processes to be started or killed as a group
 - Defines inter group dependencies

Failure Recovery

- Roll back failed operations and return an error
- Continue with reduced capacity
 - Accept requests that can be handled and reject those who can't
- Automatic restarts

- Escalation mechanisms (partial restarts, then escalating) such as restarting more groups

Making Synchronization Easier

- Identify shared resources
 - Identify what methods require serialization
- Write code to operate on those objects, and assume critical sections will be serialized
- Compiler generates the serialization automatically
 - Automatically generated locks and releases

Monitors

- Protected Classes
- Each monitor object has a mutex
 - Automatically acquired on any method invocation
 - Automatically released on method return
- Good encapsulation
 - Developers don't need to identify critical sections
 - Automatic locking when modifying a monitor object

```
monitor CheckBook {
    // object is locked when any method is invoked
    private int balance;
    public int balance() {
        return(balance);
    }
    public int debit(int amount) {
        balance -= amount;
        return( balance)
    }
}
```

Java Synchronized Methods

- Each object has an associated mutex
 - Only acquired for specific mutex
- Static synchronized methods lock class mutex
- Advantages
 - Fine grain locks and reduced deadlock risk
- Costs
 - Developer must identify serialized methods
- Object is only locked when the method attempts to modify it

METHODS

```
class CheckBook {  
    private int balance;  
    // object is not locked when this method is invoked  
    public int balance() {  
        return(balance);  
    }  
    // object is locked when this method is invoked  
    public synchronized int debit(int amount) {  
        balance -= amount;  
        return( balance)  
    }  
}
```

13