

Threads and IPC

Expensive Processes

- Creating processes takes resources
- To dispatch: freeing space
- Different processes cannot share address spaces or share resources
- Not all programs require strong separation
 - Multiple activities working together for a single goal
 - Mutually trusting elements of a system

Threads

- Strictly a unit of execution and scheduling
 - Each thread has its own stack, PC, and registers
 - Resources can be shared among threads
 - Threads are an interpreter
- Multiple threads can run a single process
 - Share the same code and data space (switching between threads doesn't take a context switch)
 - All have access to the same resources
 - Makes them cheaper to create and run
- Sharing the CPU between multiple threads
 - User level threads (using voluntary yielding)

Using Processes

- Running distinct programs
- Creating and destruction are rare events
 - (not as expensive if they are rare)
- When running agents with distinct privileges
 - (we don't want sharing)
- When there are limited interactions and shared resources
 - Each process has a distinct resource pool
- To prevent interference between executing interpreters
- Needs to isolate failures from each other

Multiple processes

- Run slowly
- Difficult to share resources

Using Threads

- Parallel activities in a single program
- When there is frequent creation and deletion
- Can be run with the same privileges
- When resource sharing is needed

Multiple threads

- Have to create and manage several threads
- Serialize resource use
- Added complexity
- Threads require protection from each other and must be managed by the programmer

Thread State and Thread Stacks

- Each thread has its own register, PS, PC
- Each thread has its own stack area
- Max stack size specified when thread is created
 - A process can contain many threads
 - Thread creator must know max required stack size (cannot grow stack as more threads added)
 - Stack space must be reclaimed when thread exits
- Procedure linkage conventions remain unchanged to processes
 - (threads can access libraries like processes)

User Level Threads vs Kernel Threads

- Kernel threads:
 - Abstraction provided by the kernel
 - Still share one address space but scheduled by the kernel
 - Multiple threads can use multiple cores
- User level threads:
 - Kernel knows nothing about them
 - Provided and managed via user level library (code in the library schedules the threads)
 - Scheduled by library, not kernel

Communications between processes

- Even fairly distinct processes may need to communicate
- OS provides mechanisms to facilitate communications
- Ex. Piping

IPC Goals

- Simplicity
- Convenience
- Generality
 - Should apply to any amount of data or signal
- Reliability
 - When data is transferred, should be certain that it works

IPC Support for OS

- Provided through sys calls
- Required activity from both communicating processes
- Usually mediated at each step by the OS
 - To protect both processes
 - Ensures correct behavior

OS IPC Mechanisms

- For local processes
- Data is in memory space of sender
- Data needs to get to memory space of receiver
- 2 choices:
 - OS copies the data
 - The OS uses VM techniques to switch ownership of memory to the receiver

Copying the data

- Conceptually simple
- Less likely to lead to user confusion
- Potentially high overhead

Using VM

- Much cheaper than copying bits
- Requires changing page tables
- Usually only one of the two processes sees the data at a time

IPC Forms

- Synchronous
 - Writes block until data is sent/delivered/received
 - Reads block until new data is available
 - Easy for programmers
- Asynchronous
 - Writes return when system accepts
 - No confirmation of transmission/delivery/reception (no error checking)
 - Reads return promptly if no data is available

Typical IPC Operations

- Create or destroy an IPC channel
- Write/send/out
 - Insert data into the channel
- Read/receive/get
 - Extract data from the channel

- Channel content query
 - Amount of data in the channel
- Connection establishment

IPC Messages vs Streams

- Stream
 - Continuous stream of bytes
 - Read or write a few or many bytes at a time
 - Write and read buffer sizes are unrelated
 - Streams may contain app specific record delimiters
- Messages
 - Sequence of distinct messages
 - Each message has its own length
 - Each message is typically read/written as a unit
 - Delivery of a message is typically everything or nothing at all

Flow Control

- **Flow Control**: ensure a fast sender doesn't overwhelm a slow receiver
- Queued IPC consumes system resources
 - Buffered in the OS until the receiver asks for it
- Fast sender or non responsive receiver can increase the amount of buffer space needed
- Must be way to limit required buffer space
 - Sender side: block or refuse connection
 - Receiving side: either tell the sender to stop/slow down or drop packets
 - Handled by network protocols or OS mechanism
- Mechanism for receiver to send feedback to sender

Reliability

- **Not a huge issue in a single machine**, OS doesn't just lose IPC data
- Across a network, requests and responses can be lost
- On a single machine, a sent message might not be processed
 - Infinite loop, buggy process, dead or invalid receiver

Pipelines

- Data flows through series of programs, like a pipe
 - Macro processor => compiler => assembler
- Data is a simple byte stream
 - Buffered in the OS
 - No need for temporary intermediate files

Sockets

- Connection between addresses and ports
 - Can connect and listen
 - One way communication
 - Sockets are complex to implement
- Many data options (flexible)
 - Reliability or best effort datagrams
 - Streams, messages
- Complex flow control and error handling
 - Retransmission and timeouts
 - Reconnection

Shared Memory

- OS arranges for processes to share read/write memory segments
 - Mapped into multiple process's address spaces
 - Application must provide their own control of sharing
 - OS not involved in the data transfer after it sets up the connection
 - Performance is very fast
- Varying complexity
 - Simple to understand
 - No synchronization, so the processes have to coordinate
- Only works on a local machine

Synchronization

- Making things happen in the right order
- Required for parallelism

Benefits of parallelism

- Improved throughput
 - Blocking on activity doesn't stop the others
- Modularity
 - Complex operations can be split
- Improved robustness
 - One thread failing doesn't stop the others

Problems with parallelism

- Sequential programs are easy to manage

- Independent parallel programs are easy if the parallel streams don't interact with each other
- Cooperating parallel programs are difficult
 - If the two execution streams aren't synchronized, results depend on the order of instructions
 - Order of execution is non deterministic

Race Conditions

- Different results occur when running threads in parallel since execution order changes every run
- Mutual exclusion: Conflicting updates
- Sleep problem: check/act races (waiting CPU should immediately start working when work is available)
- All or none transactions: multi-object updates (conflicting updates can occur)
- Distributed decisions: when there are disagreements to threads finishing

Non deterministic execution

- Parallel execution makes process behavior less predictable
 - Process block for I/O or resources
 - Time slice end preemption
 - Interrupt service routines
 - Unsynced execution on another core
 - Queuing delays
 - Time required to perform I/O operations
 - Message transmission time

Synchronization

- True parallelism is too difficult
- Pseudo parallelism is good enough
- We only need to synchronize key points of interaction

Critical section serialization

- Critical section: a resource shared by multiple interpreters
 - Sections which are shared by multiple concurrent threads, processes, or CPUs
 - Or by Interrupted code and interrupt handler
- Use of resources changes its state (ex. Write to contents or properties)
- Correctness depends on execution order
 - May not be a correct order, but it needs to be controlled

Editing and Reading files

- Updating a file by 2 processes
- One process wants to write to the file, the other wants to read

```
remove("inventory");  
fd = create("inventory");  
                                fd = open("inventory",READ);  
                                count = read(fd,buffer,length);  
write(fd,newdata,length);  
close(fd);
```

- If Process 2 reads the inventory after its been created, then it reads an empty file
- We wanted to read it after it wrote the data

Re-entrant Signals

<pre>load r1,numsigs // = 0 add r1,=1 // = 1 store r1,numsigs // =1</pre>	<pre>load r1,numsigs // = 0 add r1,=1 // = 1 store r1,numsigs // =1</pre>
---	---

- Exact same code which handles interrupts

Even A Single Instruction Can Contain a Critical Section

thread #1

thread #2

counter = counter + 1; counter = counter + 1;

*But what looks like one instruction in
C gets compiled to:*

```
mov counter, %eax  
add $0x1, %eax  
mov %eax, counter
```

- An interrupt after any instruction can cause a conflict
- We expect that the result of both running is 2
- Each has their own stack, but shares the use of registers. Saving registers onto the stack when a thread is paused
- Each thread has an inconsistent view if there is an incomplete computation left in a register on an interrupt

Critical Section and Mutual Exclusion

- Only one thread can execute a critical section at a time

Interrupt Disabled

- No interrupts: we would never interpret code working on a critical section
- Can be done with privileged instructions
- Prevents time slice ends (timer interrupts)
- Prevent re entry of device driver code

Downsides

- Delays important operations
- Bug may leave the program permanently disabled
- Won't solve all sync problems
- Can't be done in user mode

Other solutions

- Avoid sharing data
- Eliminate critical sections with atomic instructions (but difficult to only work with one instruction)
- Use locks

Summary

- Processes too expensive for soe purposes
- Threads provide a cheaper alternative
- Threads communicate through memory
- Processes need IPC