

# Intro to OS

## Operating Systems

- Performs high level complex operations
- System software to provide support for higher level applications
- Includes higher level software applications, but primarily used for user processes
- Software that sits between the hardware and everything else

## OS features

- Memory management
- Persistent storage
- Scheduling and synchronization
- Interprocess communications
- Security

## Problems that OS handles

- Coordinating separate computations
- Managing shared resources
- Virtualize hardware and software (working with different types of flash drives is done by virtualizing a device)
- **Virtualization**: OS takes a physical resource and transforms it into a more general, powerful virtual form of itself (VM)
- Organizing communications
- Protecting computing resources

## OS properties

- Services as objects
  - Represented as data structures
- Interface vs Implementation
  - Interface: interacting with the service (what you ask the service to do)
  - Implementation isn't the specification
  - Keep the implementation separate from the interface (users don't need to understand the implementation, they only use the interface)
  - Maintain the interface while changing the implementation
  - Users should avoid changing the implementation
- Interface specification
  - Acts as a contract which specifies the responsibilities of the producers and consumers
  - Basis for product release or interoperability

## OS Organization and Performance

- Modularity and function encapsulation
  - Complexity hiding and appropriate abstraction

- Each module is responsible for one thing (module for memory management)
- Simplifies the organization of an OS

- **Separate policy from mechanism**

- Policy determines what can or should be done (rules)
- Mechanism implements basic operations to do it (implementation)
- Mechanism shouldn't dictate or limit policies
- Policies must be changeable without changing the mechanism

- Parallelism

- Powerful and vital, but dangerous when used carelessly

- Performance and correctness are opposites

- Correctness doesn't always win

## OS purpose

- Low level software which provides an abstraction of hardware
- Hardware management
  - Programs interact with the hardware through the OS
  - Allocates hardware and manages its use
  - Enforces controlled sharing and privacy
  - Oversees execution and handles problems
- Hardware abstraction
  - Makes it easier to use and improves SW portability
  - Optimizes performance
- New abstractions for applications
  - OS provides new features beyond the bare hardware
  - Can install a variety of different programs

## OS structure

- A set of management and abstraction services
- OS manages computer hardware that the program isn't responsible for (turning on the fan when its too hot)
- Objects and services
  - Applications only see the objects and their services
  - CPU supports data types and operations (bytes, shorts, longs, adds, subtracts)
  - An operating system does the same at a higher level with files, processes, devices, ports, and operations like create, destroy, read, write

- Extension of a computer
  - Creating a much richer virtual computing platform

#### Running multiple programs

- Virtualizing the CPU (turning a single CPI into a seemingly infinite number)

#### OS - App Flow

1. Hardware CPU
  - Instructions divided into standard and privileged set
2. Application Software
  - Ex. games, video program
  - Can run instructions in the standard instruction set
  - Many application level instructions such as add and mov can be run directly on the hardware
3. Libraries
  - a. To interact with an application, instructions go through an Application Binding Interface (ABI)
  - b. From the library, instructions pass through the system call interface
  - c. Instructions then transferred to the OS, which translates instructions to be read by the CPU
  - This is done using the privileged instruction set (instructions that can only be done with the OS)

#### Application Binding Interface


- ABI instructions written in binary, instructions for the OS on how to instruct the CPU

#### System Call Interface

- Machine code instructions (set of operations you can ask OS to do)
- Call on the operating system to carry out instructions

#### Privileged instructions

- Instructions only the OS can carry out (app does not have permission to do so)
- Defined by CPU design (some instructions are defined as standard and others as privileged)
- Ex. Go to VM and load pointer to VM table



## OS Specialities

- Complete control over the hardware
  - Automatically loads when the computer boots
  - Continues running while apps are opened and closed
  - Privileged instruction set gives OS complete access to all memory and I/O
- Mediates access to hardware
  - OS can block, permit, or modify application requests
- Trusted
  - OS is trusted to store and manage critical data
  - Assumed to always act in good faith
- OS crash
  - If OS crashes, then everything else crashes

## Instruction Set Architecture (ISA)

- Instruction set supported by Computer
  - Which bit patterns correspond to an operation
- Many different ISAs (all incompatible)
  - Different word/bus widths (8, 16, 32 bit)
  - Different features (low power, DSPs, floating point)
  - Different design philosophies (RISC v.s CISC)
  - Competitive reasons (x86, PowerPC, Apple Silicon), companies compete by using their own unique ISA chips
- Usually come in families

- Newer models add additional features

### Privileged vs General Instructions

- Modern ISAs divide instructions into privileged vs general
- General Instructions
  - Any code running on the machine can execute general instructions
- Privileged Instructions
  - Processor must be put into specific mode to execute privileged instructions
  - OS is the only one in the special mode
  - Privileged instructions can do risky things

### Platforms

- Every computer and device is a unique platform
- ISA doesn't completely define a computer
  - Functionality beyond user mode instructions:
    - Amount of Ram or storage on a computer
    - Input output devices
- Variations are called platforms
  - The platform on which the OS must run on
  - There are lots of them

### Portability to Multiple ISAs

- If we want an OS to run on multiple different devices, it has to run on multiple ISA's
- OS must abstract the ISA
- Minimal assumptions or reliance on specific hardware
  - General frameworks are hardware independent (file systems, protocols, processes, etc)
  - Hardware assumptions isolated to specific modules (I/O, memory management)
- Still want to distribute to as many computers as possible

### Binary Distribution Model

- Binary is the derivative of source
  - The OS is written in source, and a source distribution must be compiled
  - A binary distribution is ready to run after downloading
- OSes usually distributed in binary, with one or more binary distributions per ISA
- Binary model for platform support
  - Device drivers can be added after-market

### Binary Configuration Model

- Eliminates the need for manual or static configuration

- Enables one distribution to serve all users
- Improve both ease of use and performance
- Automatic hardware discovery
  - Self identifying busses
  - We can see what is connected to each port (see details of a flashcard or keyboard)
  - **Device driver**: Each piece of hardware connected needs code on the OS to run
- Automatic resource allocation
  - OS code must make it to RAM (takes space), OS can dynamically allocated memory
  - Eliminates fixed sized resource pools
  - Dynamically allocate resources on demand
- Persistent storage
  - When a device is shut off, OS is in persistent storage
  - Moved to RAM when turned on

#### OS Functionality

- OS code is expensive, but we must include what is needed (must not have bugs to avoid crashes)
- Functionality must be in the OS if:
  - Requires the use of privileged instructions (only the OS can do this)
  - Requires the manipulation of OS data structures
  - Must maintain security, trust, or resource integrity (only the OS is trusted to maintain hardware)
  - Shared resources should be controlled by only a trustworthy source (the OS)
- Functions in library:
  - Are a service commonly needed by applications (video graphics code)
  - Doesn't need to be implemented inside OS
- Performance
  - Functions that need to be fast are written in the OS
  - If necessary functions are on a library, then permissions need to be passed back and forth, wasting time

#### Summary

- OS: software that interacts directly with the hardware
- Important for OS the remain stable