

# Mutual Exclusion and Asynchronous Completion

## Critical Section Problem

- Most common for multithreaded apps
- **Updating object state**
  - Updates to a single object
  - May be related updates to multiple objects
- **Generally involves multi step operations**
  - Object state inconsistent until operation finishes
  - Pre-emption compromises object or operation
- Correct operation requires mutual exclusion (one thread has access to the object)

## Atomicity

- **Before or After**
  - A enter critical section before B starts
  - B enters critical section after A completes
  - No overlap in entering the critical section
- **All or None**
  - An update or operation that starts must complete or will be undone
  - An uncompleted update has no effect (gets removed)
  - No incomplete computations
  - Ex. If an operation requires updating 2 data structures, then wither both get updated or non get updated
- Correctness relies on both of these

## Options for protecting critical sections

- Turn off interrupts (not practical)
- Avoid sharing data (no critical sections or only read only sections)
- Protect critical sections using hardware mutual exclusion
  - Trust CPU instructions are atomic
- Software locking

## Avoiding Shared Data

- Don't share data we don't need to share (not always practical)
- May lead to inefficient resource use
- Sharing read only data

## Atomic Instructions

- CPU instructions are uninterruptible
  - Read and modify write operations
  - Can be applied to contiguous bytes
  - Increment, decrement and binary operations

## Locking

- Protect critical sections with a data structure
- Locks
  - Party with the lock can access the critical section
  - Parties not holding the lock cant access it
- A party needing to use the critical section tries to acquire the lock
- When finished with critical section, the lock is released

## Software Locks

- ISA doesn't include instructions for building locks
- Individual instructions are serialized, while multiple instructions aren't
- Building locks in software leads to issues with mutual exclusion

## Building Locks

- The operation of locking and unlocking a lock is also a critical section
  - Must be protected to avoid 2 threads getting the same lock
- Hardware assistance
  - Individual CPU instructions are atomic
  - We can implement a lock with one instruction (either thread which runs the instruction get the lock)

## Lock Building Single Instructions

- Assembly level instruction
- Single instruction, therefore is atomic

## Test and Set

```
bool TS( char *p) {  
    bool rc;  
    rc = *p;           /* note the current value      */  
    *p = TRUE;         /* set the value to be TRUE      */  
    return rc;         /* return the value before we set it */  
}  
  
if !TS(flag) {  
    /* We have control of the critical section! */  
}
```

- Describes the testing set instructions which returns what the lock was before it was set
- Only when TS() gets called will the lock be set
- If False, we got the lock. If True, then we didn't get the lock (since the lock was True before the function call)

## Compare and Swap

```
bool compare_and_swap( int *p, int old, int new ) {
    if (*p == old) {      /* see if value has been changed
        *p = new;          /* if not, set it to new value
        return( TRUE);     /* tell caller he succeeded
    } else                 /* someone else changed *p
        return( FALSE);    /* tell caller he failed
    }

    if (compare_and_swap(flag,UNUSED,IN_USE) {
        /* I got the critical section! */
    } else {
        /* I didn't get it. */
    }
}
```

- If the lock is free, set the lock to used and return True
- Else if the lock isn't free return False

## Implementing Locks

```
bool getlock( lock *lockp ) {
    if (TS(lockp) == 0 )
        return( TRUE);
    else
        return( FALSE);
}

void freelock( lock *lockp ) {
    *lockp = 0;
}
```

## Lock Enforcement

- Locking resources works when:
  - Not possible to use a locked resource without the lock
  - Anyone who uses the resources follows locking rules
- If rules aren't followed, a thread could use the resource even when it doesn't hold the lock

## Spin Locking

- **Spin Locking**: Trying to get the lock again, it could still not be available

### Advantages

- Properly enforces access to critical sections
- Assumes properly implemented locks

### Disadvantages

- Wasteful: spinning uses processor cycles
- Likely to delay freeing of desired resource
  - Cycles burned could have been used by locking party to finish its work
- Bug could result in infinite spin wait

### Asynchronous Completion Problem

- Parallel activities move at different speeds
- Problem: One activity may need to wait for another process to complete (so it can't run in its allocated time slice)
- Async completion problem deals with how to perform waits without reducing performance
  - Waiting process goes to sleep then is woken by the OS when the blocking operation is finished
- Examples of async completion
  - Waiting for I/O operation to complete
  - Waiting for a response to a network request

### Spin Locking for Synchronization

1. When awaited operation is in parallel
  - Hardware device accepts a command
  - Another core releases a briefly held spin lock
2. When awaited operation is guaranteed to happen soon
  - Spinning is less expensive than sleeping the waking up
3. When spinning doesn't delay an awaited operation
  - Burns memory bandwidth and slows I/O
  - Burning CPU delays running another process
4. When the contention is rare
  - Multiple waiters increases the cost to run

### Yield and Spin

- Check if your event occurred
- Maybe check a few more times, then yield and allow another process to run
- Cycle of checking and yielding

### Problems

- Extra context switches
- Wastes cycles spinning each time
- Might not get scheduled to check until long after event occurs
- Works poorly with multiple waiters (could be unfair)

## Fairness and Mutual Exclusion


- Spin locking and yield and spin, but are unfair
- Both spin options require mutual exclusion

## Completion Events

- If we can't get the lock, block
- Ask the OS to wake when the lock is available
- Similar for anything else you need to wait for
  - I/O completion
  - Waiting for another process to finish
- Implemented with condition variables

## Condition Variables

- Create a synchronization object associated with a resource or request
- Requester blocks and is queued awaiting event on that object
- Posting event to object unblocks the waiter



## Condition Variables and the OS

- Generally the OS provides condition variables (library code with threads also has this)
- Block a process or thread when a condition variable is used
- It observes when the desired event occurs
- Unblocks the blocked process or thread
  - Places it back in the ready queue

## Handling Multiple Waits

- Threads can wait on several different conditions
- Pointless to wake up everyone on every event
  - Each should only wake when their event happens
- OS should allow easy selection of the right condition
- Several threads could also be waiting for the same condition

## Waiting List

- Each completion event needs an associated waiting list
- When an event occurs, consult list to determine who's waiting for that event
  - Order in which waiting list is notified depends on the event and application

## Alerting Threads

- Pthread\_cond\_wait: at least one blocked thread
- Pthread\_cond\_broadcast: all blocked threads
- Broadcast approach could be wasteful (unbound waiting times)
- Waiting queue solves these problems (a post wakes the first client in the queue)

## Locking and Waiting Lists

- Spinning is a bad thing, locks need waiting lists
- Waiting list is a shared DS
  - Will have critical sections which are protected by other locks
- Potential circular dependency

## Sleep/wakeup race condition

- Thread B has locked a resource and thread A needs to get that lock
- Thread A will call sleep() to wait for the lock to be free
- Thread B finishes using the resource
  - Thread B calls wakeup() to release the lock
- Currently no other threads waiting, just A
- However, if we never reached the code to add A to the waiting list due to a context switch for B
- Code to add A to the waiting list is reached, and A gets blocked with no one to wake up A
- The lock is free, but A is asleep

<b>Thread A</b>	<b>Thread B</b>
<pre>void sleep( eventp *e ) {     while(e-&gt;posted == FALSE) {          add_to_queue( &amp;e-&gt;queue, myproc );         myproc-&gt;runstate  = BLOCKED;         yield();     } }</pre>	<pre>void wakeup( eventp *e) {     struct proce *p;    <u>Now it happens!</u>      e-&gt;posted = TRUE;     p = get_from_queue(&amp;e-&gt; queue);     if (p) {              } /* if !p, nobody's waiting */ }</pre>

**The effect?**

Thread A is sleeping      But there's no one to  
wake him up

### Sleep() problem

- Critical section in sleep()
  - Starting before testing the posted flag
  - Ending after adding us to the waiting list and block
- During this section, we need to prevent:
  - Wakeup of the event
  - Other threads waiting on the event
- Mutual exclusion problem
  - Need a lock to control the sleep()

### Summary

- Mutual Exclusion
  - Allows only one of several things to occur at once
- Asynchronous completion
  - Properly synchronize cooperating events
- Locks are one way to assure mutual exclusion
- Spinning and completion events are ways to handle asynchronous competitions